

LLM Inference Optimization

Autumn 2025

Lecturer: Yuedong (Steven) Xu

Shenzhen Loop Area Institute

yuedongxu@slai.edu.cn

Fudan University

ydxu@fudan.edu.cn

Disclaimer


Machine learning systems is a broad and rapidly evolving field. The course material has been developed using a broad spectrum of resources, including research papers, lecture slides, blogposts, research talks, tutorial videos, and other materials shared by the research community. External pictures and are heavily reused.

Inference Optimization: Outline

- Overview
- Attention Optimization
- Continuous Batching
- KV Cache Optimization
- Speculative Decoding
- Distributed Serving

Overview


- Decoder-only Transformer



GPT (Generative Pre-trained Transformer) is the first decoder-only Transformer model

Overview


- Decoder-only Transformer
 - Generating a probabilistic distribution over possible next token, and a decoding algorithm is employed to select the actual output token



Next-token prediction, i.e. generating output tokens one by one


Overview

- Autoregressive Decoding
 - **“Prefill”** refers to the **initial parallel computation phase** where the model processes the entire input prompt for subsequent token-by-token generation




Overview

- Prefill: highly parallelizable
 - A small batch size can "saturate" GPU computation
 - Parallelization over batch size, header size, sequence length and thread-block tiling




Overview

- Autoregressive Decoding
 - “**Decode**” refers to the **iterative process of generating output tokens** one at a time, where each new token is predicted based on the input prompt and all previously generated tokens

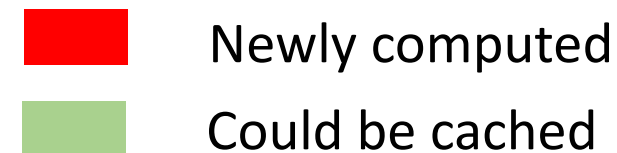


Overview


- Decode: much less computes



Overview




- Decode: much less computes
 - Why caching Key and Value, other than Query



Decode Compute output token one by one

Overview

- An animation



Overview


- Why Caching KV?
 - Reducing redundant computations for KEYS and VALUES
 - Increasing memory consumption

2 × 2 × 8096 × 80 × 64 × 64
K/V Float16 **Sequence length** # of layers # of heads dimension

~10.6 GB!


Overview

- Autoregressive Decoding
 - Token generation core metrics: **TTFT** and **TPOT** (time per output token)



Overview


- Hierarchical GPU Memory
 - I/O throughput versus memory size



Memory Hierarchy with
Bandwidth & Memory Size

A macroscopic view of I/O tradeoff (V100)


Insufficient to store data at
SRAM: load from HBM to
SRAM and write back to HBM



L1 instruction cache: 192KB per SM * 108 SM ~20MB
Data flow: DRAM/L2 Cache to/from L1 Cache, much
slower than computing

Overview


- Prefill and Decode
 - Prefill: Compute intensive with GEMM, Decode: memory I/O intensive with GEMV



Roofline model of NVIDIA V100 GPU

Overview

- Prefill and Decode
 - Prefill: Compute intensive, Decode: memory I/O intensive
 - Prefill saturates GPU compute even at batch size of 1
 - Decode under-utilizes GPU compute and costs as much as 200 times prefill for bs=1



Per-token prefill and decode time with different batch sizes
(sequence length = **1024**) for LLaMa-13B on A6000 GPU


Overview

- Very Large-scale LLM Inference




Full model size with 256 experts

Long context and high
request loads




Overview

- Versatile communication medium
 - Transmission delay emerges



PCIe Link
e.g. CPU – GPU with
up to 128GB/s bandwidth



NVLink
e.g. GPU – GPU
in total 1.8TB/s bandwidth



RoCE
e.g. Machine – Machine
up to 800 Gbps


Overview

- Challenges of LLM Inference
 - New token generation paradigm
 - **Prefill** and **Decode**
 - Hierarchical memory
 - **Fast** I/O small size versus **slow** I/O large size
 - Heterogeneous bandwidth
 - **High** intra-machine bandwidth versus **low** inter-machine bandwidth
- To emphasize
 - many optimization methods, e.g. **attention optimization** and **prefill optimization** can be employed for model **training**

Inference Optimization: Outline

- Overview
- Attention Optimization
 - **Sparse Attention**
 - Linear Attention
 - Flash Attention
 - Continuous Batching
- Continuous Batching
- KV Cache Optimization
- Speculative Decoding
- Distributed Serving

Attention Optimization



Attention Optimization

- Real-world operations
 - Transferring both matrices from global memory to shared memory for computing, and write the result back to global memory progressively
 - Complexity
 - Computation: $O(N^2d)$
 - Considering multiplications and additions $\rightarrow$ computing time $T_C = \frac{2N^2d}{c}$
 - Communication: $O(N^2)$
 - Considering bidirectional I/O read/write $\rightarrow$ communication time $T_{I/O} = \frac{2*(N^2+2Nd)}{B}$
 - Global storage: $O(N^2)$
-
- (partially) overlapped**
- GPT3 with 10K tokens in FP16:
200MB per head per layer

Attention Optimization

- Real-world operations
 - Transferring both matrices from global memory to shared memory for

Computation, I/O and storage should all be optimized!


- Communication: $O(N^2)$
 - Considering bidirectional I/O read/write $\rightarrow$ communication time $T_{I/O} = \frac{2*(N^2+2Nd)}{B}$
- Global storage: $O(N^2)$

Sparse Attention

- Sparse attention
 - reduce computational and memory cost by only computing attention for a **subset of token pairs** instead of all pairs
 - evidence from machine learning community (e.g. Rewon Child, 2019)
- Sparse attention patterns
 - Position-based sparse attention: **rule-based or heuristic methods** to decide the positions where the interactions of these pair-wise tokens are important
 - Content-based sparse attention: tokens selectively attending only to other tokens that are **relevant based on their representations**

Position-based Sparse Attention

- Global attention
 - global nodes act as an information hub, allowing them to attend to every other node in the sequence



A bipartite graph illustration of $\mathbf{QK}^T$ computation

Key

k_1 k_2 k_3 k_4 k_5 k_6 k_7 k_8

Query

q_1

q_2

q_3

q_4

q_5

q_6


q_7

q_8

Not stored in memory

Position-based Sparse Attention

- Band attention
 - often termed “local attention” or “sliding window attention”, in which a node’s attentions are confined to neighboring nodes in a local window




	Key							
	k_1	k_2	k_3	k_4	k_5	k_6	k_7	k_8
Query	q_1							
	q_2							
	q_3							
	q_4							
	q_5							
	q_6							
	q_7							
	q_8							

- Reducing complexity from $O(N^2)$ to $O(kN)$
- Limited receptive field (cannot capture long-range dependencies)
- Sacrificing context modeling ability: slow information propagation

Position-based Sparse Attention

- Dilated attention
 - using a dilated window with gaps of dilation equal to or greater than 1





	Key							
	k_1	k_2	k_3	k_4	k_5	k_6	k_7	k_8
Query q_1								
q_2								
q_3								
q_4								
q_5								
q_6								
q_7								
q_8								

- Reducing complexity from $O(N^2)$ to $O(kN)$
- Missing fine-grained, short-range dependencies (failure to capture important neighbors in those gaps)


Position-based Sparse Attention

- Random attention
 - each query randomly samples a few keys for better capturing non-local interactions



Position-based Sparse Attention


- Block attention
 - input sequence segmented into multiple non-intersection query blocks and each block assigned a local memory block




	Key							
	k_1	k_2	k_3	k_4	k_5	k_6	k_7	k_8
Query	q_1							
	q_2							
	q_3							
	q_4							
	q_5							
	q_6							
	q_7							
	q_8							

Position-based Sparse Attention


- COMPOUND attention




(a) Star-Transformer



(b) Longformer



(c) ETC




(d) BigBird

- BigBird : local + global + random yields dense-equivalent performance
- Increased implementation complexity, more hyperparameters to tune and potential redundancy

Position-based Sparse Attention


- Representative sparse attentions in a nutshell



Content-based Sparse Attention


- Routing transformer

Clustering (approximate & fast)




Content-based Sparse Attention

- Routing transformer




K-means Clustering (approximate & fast)



Content-based Sparse Attention

- Clusters

- $\{q_1, q_4, q_6, k_1, k_3\}$
- $\{q_2, k_1, k_4, k_8\}$
- $\{q_3, q_5, k_6\}$
- $\{q_7, q_8, k_2, k_5, k_7\}$




Content-based Sparse Attention

- Routing transformer: summary
 - attention scores in scaled dot-product rely on the similarity between Q and K
→ approximated by grouping them in a shared projection space
 - needs to be **double checked** by our students manually
 - each query only needs to attend to keys within its assigned cluster (typically $\sqrt{n}$ in size), reducing complexity to $O(n\sqrt{n})$ while preserving relevant long-range dependencies
 - cluster centroids are learned during training, allowing the model to adaptively group based on content

Content-based Sparse Attention


- Reformer: using locality-sensitive hashing (LSH) instead of dot-product attention




- Using **locality-sensitive hashing** (LSH) to select key-value pairs for each query
- Using Reversible Transformer to reduce memory usage during training
- Similar items (queries and keys) falling in the same bucket with high probability
- Processing sequences of length **64K tokens** or more efficiently

Inference Optimization: Outline


- Overview
- Attention Optimization
 - Sparse Attention
 - **Linear Attention**
 - Flash Attention
- Continuous Batching
- KV Cache Optimization
- Speculative Decoding
- Distributed Serving



“optimal parenthesization” via dynamic programming





$$O(N d_k N)$$



$$O(N N d_v)$$

Original


$$O(N^2(d_v + d_k))$$



$$O(N N d_v)$$


Original

$$O(N^2(d_v + d_k))$$




Original

$$O(N^2(d_v + d_k))$$




New

$$O(N d_k d_v)$$



Letting $N = 2048$
 $d_k = d_v = 64$

Original



$$O(N^2(d_v + d_k)) = 536,870,912$$


$$O(d_k N d_v)$$



1.56%!

New

$$O(N d_k d_v) = 8,388,608$$



$$O(N d_k d_v)$$

Obstacle

- Softmax prevents this reordered GEMM

$$\text{Softmax} \left(\frac{1}{\sqrt{d_k}} \begin{matrix} d_k \\ N \\ Q \end{matrix} \begin{matrix} N \\ d_k \\ K^T \end{matrix} \right) \begin{matrix} d_v \\ N \\ V \end{matrix}$$

Linear Attention

- Computing Softmax (Katharopoulos et al., 2020)

$$o_1 = \sum_{i=1}^N \alpha_{1,i} v_i = \sum_{i=1}^N \frac{\exp(q_1 \cdot k_i)}{\sum_{j=1}^N \exp(q_1 \cdot k_j)} v_i$$

- Substituting $\exp$ by

$$\exp(q \cdot k) \approx \phi(q) \cdot \phi(k)$$

- there exists

$$o_1 \approx \sum_{i=1}^N \frac{\phi(q_1) \cdot \phi(k_i)}{\sum_{j=1}^N \phi(q_1) \cdot \phi(k_j)} v_i$$


Kernel function

Linear Attention

- Overall Output: complexity order $O(N)$

$$O \approx \phi(Q) \cdot (\phi(K) V), \quad O, Q, K, V \in \mathbb{R}^{N \times d}$$

- An example: denominator of o_1



Linear Attention

- Iterative computation to further reduce complexity

$$o_1 \approx \sum_{i=1}^N \frac{\phi(q_1) \cdot \phi(k_i)}{\sum_{j=1}^N \phi(q_1) \cdot \phi(k_j)} v_i$$

- where

$$\sum_{i=1}^N [\phi(q_1) \cdot \phi(k_i)] v_i = \phi(q_1)^\top \left(\sum_{i=1}^N k_i v_i^\top \right) \Rightarrow$$

$$S_t = \sum_{i=1}^t k_i v_i^\top \in R^{d \times d}$$

$$S_t = S_{t-1} + k_t v_t^\top$$

$$o_t = \phi(q_t)^\top S_t$$

Linear Attention

- Today's Linear Attention
 - Baby linear attention: Linear Transformer, SANA, CHELA, LightningAttention, etc.
 - **Efficient** in computation, I/O, storage
 - Performance loss compared with ***Softmax***
 - More advanced linear attention:
 - Delta rule
 - Gamma forget
 - Gated attention
 - Qwen3-Next and Kimi-linear, and many others

Inference Optimization: Outline

- Overview
- Attention Computation Optimization
 - Sparse Attention
 - Linear Attention
 - **Flash Attention**
 - Continuous Batching
- KV Cache Optimization
- Speculative Decoding
- Distributed Serving

Flash Attention

- Prior works
 - Reducing # of scaled dot-product, potentially sacrificing attention performance
 - Compute is fast while reading K and V , and writing S back the results are slow
- Challenges
 - I/O cost of loading/storing Q , K and V are non-trivial

Algorithm 0 Standard Attention Implementation

Require: Matrices $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d}$ in HBM.

- 1: Load $\mathbf{Q}, \mathbf{K}$ by blocks from HBM, compute $\mathbf{S} = \mathbf{QK}^\top$, write $\mathbf{S}$ to HBM.
 - 2: Read $\mathbf{S}$ from HBM, compute $\mathbf{P} = \text{softmax}(\mathbf{S})$, write $\mathbf{P}$ to HBM.
 - 3: Load $\mathbf{P}$ and $\mathbf{V}$ by blocks from HBM, compute $\mathbf{O} = \mathbf{PV}$, write $\mathbf{O}$ to HBM.
 - 4: Return $\mathbf{O}$.
-

Flash Attention

- From Softmax to Safe Softmax

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}} \quad \Rightarrow \quad \text{softmax}(x_i) = \frac{e^{x_i - \max(x)}}{\sum_{j=1}^n e^{x_j - \max(x)}}$$

- Online Safe Softmax: not enough space to store entire x

- Alg1: three-pass algorithm

m_i : $\max_{j=1}^i \{x_j\}$ with initial value $m_0 = -\infty$

d_i : $\sum_{j=1}^i e^{x_j - m_N}$ is the denominator of Softmax

a_i : is the final Softmax output

for $i \leftarrow 1, N$ do

$$m_i \leftarrow \max(m_{i-1}, x_i)$$

end

for $i \leftarrow 1, N$ do

$$d_i \leftarrow d_{i-1} + e^{x_i - m_N}$$

end

for $i \leftarrow 1, N$ do

$$a_i \leftarrow \frac{e^{x_i - m_N}}{d_N}$$

end

Flash Attention

- Online Safe Softmax
 - d_i can be computed iteratively

$$\begin{aligned} d'_i &= \sum_{j=1}^i e^{x_j - m_i} \\ &= \left(\sum_{j=1}^{i-1} e^{x_j - m_i} \right) + e^{x_i - m_i} \\ &= \left(\sum_{j=1}^{i-1} e^{x_j - m_{i-1}} \right) e^{m_{i-1} - m_i} + e^{x_i - m_i} \\ &= d'_{i-1} e^{m_{i-1} - m_i} + e^{x_i - m_i} \end{aligned}$$



Algorithm 2-pass online softmax

for $i \leftarrow 1, N$ do

end

for $i \leftarrow 1, N$ do

end

$$m_i \leftarrow \max(m_{i-1}, x_i)$$

$$d'_i \leftarrow d'_{i-1} e^{m_{i-1} - m_i} + e^{x_i - m_i}$$

$$a_i \leftarrow \frac{e^{x_i - m_N}}{d'_N}$$

Flash Attention

- Online Safe Softmax
 - One-pass safe softmax algorithm does not exist
 - Safe softmax is only our intermediate result:

$$O = \text{Softmax}(QK^T)V$$

NOTATIONS

$Q[k,:]$: the k -th row vector of Q matrix.

$K^T[:,i]$: the i -th column vector of K^T matrix.

$O[k,:]$: the k -th row of output O matrix.

$V[i,:]$: the i -th row of V matrix.

$\{\mathbf{o}_i\}$: $\sum_{j=1}^i a_j V[j,:]$, a row vector storing partial aggregation result $A[k,:i] \times V[:,i]$

BODY

for $i \leftarrow 1, N$ **do**

$$x_i \leftarrow Q[k,:] K^T[:,i]$$

$$m_i \leftarrow \max(m_{i-1}, x_i)$$

$$d'_i \leftarrow d'_{i-1} e^{m_{i-1} - m_i} + e^{x_i - m_i}$$

end

for $i \leftarrow 1, N$ **do**

$$a_i \leftarrow \frac{e^{x_i - m_N}}{d'_N}$$

$$\mathbf{o}_i \leftarrow \mathbf{o}_{i-1} + a_i V[i,:]$$

end

$$O[k,:] \leftarrow \mathbf{o}_N$$

Flash Attention

- One-pass Attention

- Define an “intermediate” output

$$\mathbf{o}'_i := \left(\sum_{j=1}^i \frac{e^{x_j - m_i}}{d'_i} V[j, :] \right)$$

for $i \rightarrow 1, N$

$$x_i \leftarrow Q[k, :] K^T[:, i]$$

$$m_i \leftarrow \max(m_{i-1}, x_i)$$

$$d'_i \leftarrow d'_{i-1} e^{m_{i-1} - m_i} + e^{x_i - m_i}$$

$$\mathbf{o}'_i = \mathbf{o}'_{i-1} \frac{d'_{i-1} e^{m_{i-1} - m_i}}{d'_i} + \frac{e^{x_i - m_i}}{d'_i} V[i, :]$$

end

$$O[k, :] \leftarrow \mathbf{o}'_N$$

Flash Attention

- Compute it iteratively

$$\mathbf{o}'_i = \sum_{j=1}^i \frac{e^{x_j - m_i}}{d'_i} V[j, :]$$

$$= \left(\sum_{j=1}^{i-1} \frac{e^{x_j - m_i}}{d'_i} V[j, :] \right) + \frac{e^{x_i - m_i}}{d'_i} V[i, :]$$

$$= \left(\sum_{j=1}^{i-1} \frac{e^{x_j - m_{i-1}}}{\underline{d'_{i-1}}} \frac{e^{x_j - m_i}}{\underline{e^{x_j - m_{i-1}}}} \frac{\underline{d'_{i-1}}}{d'_i} V[j, :] \right) + \frac{e^{x_i - m_i}}{d'_i} V[i, :]$$

$$= \left(\sum_{j=1}^{i-1} \frac{e^{x_j - m_{i-1}}}{d'_{i-1}} V[j, :] \right) \frac{d'_{i-1}}{d'_i} e^{m_{i-1} - m_i} + \frac{e^{x_i - m_i}}{d'_i} V[i, :]$$

$$= \mathbf{o}'_{i-1} \frac{d'_{i-1} e^{m_{i-1} - m_i}}{d'_i} + \frac{e^{x_i - m_i}}{d'_i} V[i, :]$$

Flash Attention


• Pseudocode of Flash Attention


Tiling: loading and computing at the block granularity


Algorithm 1 FLASHATTENTION


Require: Matrices $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d}$ in HBM, on-chip SRAM of size M .


- 1: Set block sizes $B_c = \lceil \frac{M}{4d} \rceil, B_r = \min(\lceil \frac{M}{4d} \rceil, d)$.
 - 2: Initialize $\mathbf{O} = (0)_{N \times d} \in \mathbb{R}^{N \times d}, \ell = (0)_N \in \mathbb{R}^N, m = (-\infty)_N \in \mathbb{R}^N$ in HBM.
 - 3: Divide $\mathbf{Q}$ into $T_r = \lceil \frac{N}{B_r} \rceil$ blocks $\mathbf{Q}_1, \dots, \mathbf{Q}_{T_r}$ of size $B_r \times d$ each, and divide $\mathbf{K}, \mathbf{V}$ into $T_c = \lceil \frac{N}{B_c} \rceil$ blocks $\mathbf{K}_1, \dots, \mathbf{K}_{T_c}$ and $\mathbf{V}_1, \dots, \mathbf{V}_{T_c}$, of size $B_c \times d$ each.
 - 4: Divide $\mathbf{O}$ into T_r blocks $\mathbf{O}_1, \dots, \mathbf{O}_{T_r}$ of size $B_r \times d$ each, divide ℓ into T_r blocks $\ell_1, \dots, \ell_{T_r}$ of size B_r each, divide m into T_r blocks $m_1, \dots, m_{T_r}$ of size B_r each.
 - 5: **for** $1 \leq j \leq T_c$ **do** Outer Loop on Key/Value
 - 6: Load $\mathbf{K}_j, \mathbf{V}_j$ from HBM to on-chip SRAM.
 - 7: **for** $1 \leq i \leq T_r$ **do** Inner Loop on Key/Value
 - 8: Load $\mathbf{Q}_i, \mathbf{O}_i, \ell_i, m_i$ from HBM to on-chip SRAM.
 - 9: On chip, compute $\mathbf{S}_{ij} = \mathbf{Q}_i \mathbf{K}_j^T \in \mathbb{R}^{B_r \times B_c}$.
 - 10: On chip, compute $\tilde{m}_{ij} = \text{rowmax}(\mathbf{S}_{ij}) \in \mathbb{R}^{B_r}, \tilde{\mathbf{P}}_{ij} = \exp(\mathbf{S}_{ij} - \tilde{m}_{ij}) \in \mathbb{R}^{B_r \times B_c}$ (pointwise), $\tilde{\ell}_{ij} = \text{rowsum}(\tilde{\mathbf{P}}_{ij}) \in \mathbb{R}^{B_r}$.
 - 11: On chip, compute $m_i^{\text{new}} = \max(m_i, \tilde{m}_{ij}) \in \mathbb{R}^{B_r}, \ell_i^{\text{new}} = e^{m_i - m_i^{\text{new}}} \ell_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\ell}_{ij} \in \mathbb{R}^{B_r}$.
 - 12: Write $\mathbf{O}_i \leftarrow \text{diag}(\ell_i^{\text{new}})^{-1} (\text{diag}(\ell_i) e^{m_i - m_i^{\text{new}}} \mathbf{O}_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\mathbf{P}}_{ij} \mathbf{V}_j)$ to HBM.
 - 13: Write $\ell_i \leftarrow \ell_i^{\text{new}}, m_i \leftarrow m_i^{\text{new}}$ to HBM.
 - 14: **end for**
 - 15: **end for**
 - 16: Return $\mathbf{O}$.
-




































Flash Attention

- Flash Attention (Tiling)




Flash Attention


- Kernel Fusion

- Multiple separate operations or kernels (small, independent programs that run on hardware like GPUs) are combined or "fused" into a single, more efficient kernel.



Kernel 1: Scale



Kernel 2: ReLU



Kernel 3: Add




Pay the launch overhead **once**
instead of many times


Flash Attention

- Kernel Fusion

- Minimizing data movement between slow global memory and faster on-chip memory (like SRAM or registers), avoids storing intermediate results



Multiple smaller memory transactions




Single memory transaction for 32 consecutive elements

Flash Attention

- Kernel Fusion

- FlashAttention kernel fusion combining multiple (block-wise) steps of the attention computation—such as scaling, masking, softmax normalization, and matrix multiplications—into custom fused CUDA kernels



6 - 7 × improvement

Flash Attention

- Complexity Analysis


- GPT2 Medium (seq. length 1024, head dim. 64, 16 heads, batch size 64) on A100

Attention	Standard	FLASHATTENTION
GFLOPs	66.6	75.2
HBM R/W (GB)	40.3	4.4
Runtime (ms)	41.7	7.3

- Computational complexity: Flash Attention is slightly higher
- Memory complexity: Flash Attention is significantly better $O(Nd)$
- I/O complexity
 - Standard Attention $O(N^2d)$ versus Flash Attention $O(N^2d^2M^{-1})$
 - Outer loop runs $T_c = N/B_c = 4d/M$ where M is SRAM size of a SM, and reads K and V once
 - Inner loop reads Q , O , and write O back to HBM
 - Total I/O complexity: $O(Nd) \times T_c = O(N^2d^2M^{-1})$


Flash Attention

- Flash Attention in practice
 - Impact of hyper parameters:
Larger block size reduces the number of data loading, thus improving the forward runtime
- As hidden dimension becomes large, the block size shrinks
 - $O(N^2 d^2 M^{-1})$ is quadratic to the hidden dimension d , while M is constrained by the hardware



Flash Attention

- Flash Attention
 - Underutilized streaming multi-processors for **small batch size** but **long sequences**




Grid of Thread Blocks

- GPU grid
 - Flash Attention enables Parallel computing on **batches and heads**
 - The grid consists of $batch_size \times head_number$ thread blocks
 - Batch_size = 2 and head_number = 16, meaning that only 16 out of 108 SMs are used on A100


Flash Attention

- Flash Attention v2
 - Making FlashAttention **sequence parallel**
 - The shape of a grid: `grid(num_m_block, batch_size, num_heads)`



Flash Attention

- Flash Attention v2: **Loop Reordering**



$$m^{(1)} = \text{rowmax}(\mathbf{S}^{(1)}) \in \mathbb{R}^{B_r}$$

$$\ell^{(1)} = \text{rowsum}(e^{\mathbf{S}^{(1)} - m^{(1)}}) \in \mathbb{R}^{B_r}$$

$$\mathbf{O}^{(1)} = e^{\mathbf{S}^{(1)} - m^{(1)}} \mathbf{V}^{(1)} \in \mathbb{R}^{B_r \times d}$$

$$m^{(2)} = \max(m^{(1)}, \text{rowmax}(\mathbf{S}^{(2)})) = m$$

$$\ell^{(2)} = e^{m^{(1)} - m^{(2)}} \ell^{(1)} + \text{rowsum}(e^{\mathbf{S}^{(2)} - m^{(2)}}) = \text{rowsum}(e^{\mathbf{S}^{(1)} - m}) + \text{rowsum}(e^{\mathbf{S}^{(2)} - m}) = \ell$$

$$\tilde{\mathbf{P}}^{(2)} = \text{diag}(\ell^{(2)})^{-1} e^{\mathbf{S}^{(2)} - m^{(2)}}$$

$$\tilde{\mathbf{O}}^{(2)} = \text{diag}(e^{m^{(1)} - m^{(2)}})^{-1} \tilde{\mathbf{O}}^{(1)} + e^{\mathbf{S}^{(2)} - m^{(2)}} \mathbf{V}^{(2)} = e^{s^{(1)} - m} \mathbf{V}^{(1)} + e^{s^{(2)} - m} \mathbf{V}^{(2)}$$

$$\mathbf{O}^{(2)} = \text{diag}(\ell^{(2)})^{-1} \tilde{\mathbf{O}}^{(2)} = \mathbf{O}.$$

FlashAttention forward pass: key K is partitioned into two blocks and value V is also partitioned into two blocks

Flash Attention

- FlashAttention v2


- Algorithm: fewer *non-matmul* FLOPs

- GPU is highly optimized on *matmul* (e.g. Tensor Cores)

- NVIDIA A100 GPU: 312 TFLOPs/s of FP16/BF16 *matmul*, only 19.5 TFLOPs/s of *non-matmul* FP32

$$\begin{aligned}
 S_{ij} &\leftarrow Q_i K_j^T \\
 m_{ij} &\leftarrow \text{rowmax}(S_{ij}) \\
 m_i^{new} &\leftarrow \max(m_i, m_{ij}) \\
 P_{ij} &\leftarrow \exp(S_{ij} - m_i^{new}) \\
 l_{ij} &\leftarrow \text{rowsum}(P_{ij}) \\
 l_i^{new} &\leftarrow \exp(m_i - m_i^{new}) l_i + l_{ij} \\
 O_i^{new} &\leftarrow \underbrace{(l_i^{new})^{-1}}_{\text{Scaling at this round in FlashAttention v1}} \underbrace{(l_i \exp(m_i - m_i^{new}) O_i + P_{ij} V_j)}_{\text{Scaling at next round}}
 \end{aligned}$$


Scaling at this round in FlashAttention v1





$$\begin{aligned}
 \tilde{O}_i^{new} &\leftarrow \exp(m_i - m_i^{new}) \tilde{O}_i + P_{ij} V_j \\
 O_i^{last} &\leftarrow \underbrace{(l_i^{last})^{-1}}_{\uparrow} \tilde{O}_i^{last}
 \end{aligned}$$


Only scaling the final result in FlashAttention v2


FlashAttention v2: reduce the number of rescaling ops, as well as bound-checking and causal masking operations



































Flash Attention

Algorithm 1 FLASHATTENTION-2 forward pass

Require: Matrices $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d}$ in HBM, block sizes B_c, B_r .

- 1: Divide $\mathbf{Q}$ into $T_r = \left\lceil \frac{N}{B_r} \right\rceil$ blocks $\mathbf{Q}_1, \dots, \mathbf{Q}_{T_r}$ of size $B_r \times d$ each, and divide $\mathbf{K}, \mathbf{V}$ into $T_c = \left\lceil \frac{N}{B_c} \right\rceil$ blocks $\mathbf{K}_1, \dots, \mathbf{K}_{T_c}$ and $\mathbf{V}_1, \dots, \mathbf{V}_{T_c}$, of size $B_c \times d$ each.
 - 2: Divide the output $\mathbf{O} \in \mathbb{R}^{N \times d}$ into T_r blocks $\mathbf{O}_1, \dots, \mathbf{O}_{T_r}$ of size $B_r \times d$ each, and divide the logsumexp L into T_r blocks $L_1, \dots, L_{T_r}$ of size B_r each.
 - 3: **for** $1 \leq i \leq T_r$ **do**
 - 4: Load $\mathbf{Q}_i$ from HBM to on-chip SRAM.
 - 5: On chip, initialize $\mathbf{O}_i^{(0)} = (0)_{B_r \times d} \in \mathbb{R}^{B_r \times d}$, $\ell_i^{(0)} = (0)_{B_r} \in \mathbb{R}^{B_r}$, $m_i^{(0)} = (-\infty)_{B_r} \in \mathbb{R}^{B_r}$.
 - 6: **for** $1 \leq j \leq T_c$ **do**
 - 7: Load $\mathbf{K}_j, \mathbf{V}_j$ from HBM to on-chip SRAM.
 - 8: On chip, compute $\mathbf{S}_i^{(j)} = \mathbf{Q}_i \mathbf{K}_j^T \in \mathbb{R}^{B_r \times B_c}$.
 - 9: On chip, compute $m_i^{(j)} = \max(m_i^{(j-1)}, \text{rowmax}(\mathbf{S}_i^{(j)})) \in \mathbb{R}^{B_r}$, $\tilde{\mathbf{P}}_i^{(j)} = \exp(\mathbf{S}_i^{(j)} - m_i^{(j)}) \in \mathbb{R}^{B_r \times B_c}$ (pointwise), $\ell_i^{(j)} = e^{m_i^{(j-1)} - m_i^{(j)}} \ell_i^{(j-1)} + \text{rowsum}(\tilde{\mathbf{P}}_i^{(j)}) \in \mathbb{R}^{B_r}$.
 - 10: On chip, compute $\mathbf{O}_i^{(j)} = \text{diag}(e^{m_i^{(j-1)} - m_i^{(j)}})^{-1} \mathbf{O}_i^{(j-1)} + \tilde{\mathbf{P}}_i^{(j)} \mathbf{V}_j$.
 - 11: **end for**
 - 12: On chip, compute $\mathbf{O}_i = \text{diag}(\ell_i^{(T_c)})^{-1} \mathbf{O}_i^{(T_c)}$.
 - 13: On chip, compute $L_i = m_i^{(T_c)} + \log(\ell_i^{(T_c)})$.
 - 14: Write $\mathbf{O}_i$ to HBM as the i -th block of $\mathbf{O}$.
 - 15: Write L_i to HBM as the i -th block of L .
 - 16: **end for**
 - 17: Return the output $\mathbf{O}$ and the logsumexp L .
-

Flash Attention v2: Loop Reordering


- Maintaining the same order of I/O complexity: $O(N^2 d^2 M^{-1})$
- Reducing read/write of $\mathbf{O}_i$: write $\mathbf{O}_i$ to HBM after traversing all K/V blocks at the inner loop

Flash Attention


- FlashAttention v2

- Better Work Partitioning

- A thread block is partitioned into multiple warps (32 threads per warp, 4~8 warps per block)




- Slicing K/V at each Warp while maintaining complete Q
- QK^T has four partitions, and $(QK^T)V$ needs REDUCTION (write to shared memory, synchronize, and add-up intermediate results)



- Slicing Q while keeping K/V accessible by all Warps
- Each warp performs matrix multiply to obtain a slice of QK^T , and no intra-block communication


Flash Attention

- FlashAttention v2
 - Better Work Partitioning
 - A thread block is partitioned into multiple warps (32 threads per warp, 4~8 warps per block)




Flash Attention

- FlashAttention v2
 - Performance speedup



(a) Without causal mask, head dimension 64




(b) Without causal mask, head dimension 128


Attention forward + backward speed on A100 GPU

Flash Attention

- FlashAttention v2
 - Performance speedup



(a) Without causal mask, head dimension 64




(b) Without causal mask, head dimension 128

Attention forward + backward speed on H100 GPU

Flash Attention

- Flash Decoding
 - New challenges in autoregressive decoding

Output



Input




Query (Q) length is always 1, i.e. the newly generated token

- Key (K) and Value (V) keep growing to be extraordinarily large
 - Out of Memory error
 - Swapping to CPU memory which is very slow
 - Truncating the input causes poor performance


Flash Attention

- Flash Decoding
 - Low efficiency due to small Q (i.e. sequence length of 1 in decoding)



Flash Attention

- Flash Decoding
 - Sharding Key and Value matrices (e.g. 5 pieces) in order to create many more thread blocks




Flash Attention

- Flash Decoding
 - Computing attention scores with K/V splits using FlashAttention v1/v2
 - Only obtaining local ***max*** and local ***sum***
 - Output O not properly scaled
 - Launching an independent ***REDUCE*** kernel
 - Obtaining global ***max*** and global ***sum***
 - Rescaling local O to obtain the final output

Flash Decoding

- Flash Decoding




Setting \ Algorithm	PyTorch Eager	Flash-Attention v2.0.9	Flash-Decoding
B=256, seqlen=256	3058.6	390.5	63.4
B=128, seqlen=512	3151.4	366.3	67.7
B=64, seqlen=1024	3160.4	364.8	77.7
B=32, seqlen=2048	3158.3	352	58.5
B=16, seqlen=4096	3157	401.7	57
B=8, seqlen=8192	3173.1	529.2	56.4
B=4, seqlen=16384	3223	582.7	58.2
B=2, seqlen=32768	3224.1	1156.1	60.3
B=1, seqlen=65536	1335.6	2300.6	64.4
B=1, seqlen=131072	2664	4592.2	106.6


Flash Decoding versus Flash Attention

Flash Attention v3

- Hardware improvement (only for your reference)
 - TMA (Tensor Memory Accelerator) Warpgroup Matrix Multiply-Accumulate



- Asynchrony
 - Overlap overall computation and data movement via warp-specialization and interleave block-wise **matmul** and **softmax** operations




Inference Optimization: Outline

- Overview
- Attention Computation Optimization
 - Sparse Attention
 - Linear Attention
 - Flash Attention
- **Continuous Batching**
- KV Cache Optimization
- Speculative Decoding
- Distributed Serving

Batching to Meet Compute-Bound


- Why batching multiple requests?
 - Generating tokens for a large number of prompts to match the compute bound




Roofline model of NVIDIA V100 GPU

Static Batching

- Naively serving multiple requests simultaneously
 - Unequal input sequence lengths, unknown output sequence lengths




Static Batching




Dynamic Batching

- Insert an inference request in the queue




Idea Batching

- Concatenating prompts
 - **Decode** is equivalent to **prefill** with sequence length of 1
 - **Decode** and **Prefill** of different prompts operate simultaneously without bubbles




Continuous Batching

- Continuous batching = ragged batching + ideal dynamic batching
 - Attention scores of tokens belonging to *prompt 0* and *prompt 1* should not be computed together
 - Ragged batching because sequence lengths are 'ragged' or uneven, no need for padding tokens



Continuous Batching


- Question remains unsolved: how to make **prefill** of new requests and **decode** of existing requests operate *in parallel*?



Continuous batching without chunked prefill

Chunked Prefill

- **Chunked Prefill**
 - **Elastic Scheduling:** Instead of processing the entire prompt at once, one can split it into smaller chunks: more scheduling flexibility and enables mixing prefill and decode.
 - **Decode-Maximal Batching:** one can combine a single prefill chunk with multiple decode requests in the same batch, allowing decode operations to "piggyback" on the more compute-intensive prefill operations



Chunked Prefill

- Combined together to acquire a global view of complete prefill

	k0	k1	k2	k3
q0	1	-	-	-
q1	1	1	-	-
q2	1	1	1	-
q3	1	1	1	1

attention mask during first chunk prefill

Mask is a lower triangular matrix

K/V being loaded 2 times
by late

	k0	k1	k2	k3	k4	k5	k6	k7
q4	1	1	1	1	1	-	-	-
q5	1	1	1	1	1	1	-	-
q6	1	1	1	1	1	1	1	-
q7	1	1	1	1	1	1	1	1

attention mask during second chunk prefill

Mask is a trapezoid matrix

K/V being loaded 1 time

	k0	k1	k2	k3	k4	k5	k6	k7	k8	k9	k10	k11
q8	1	1	1	1	1	1	1	1	1	-	-	-
q9	1	1	1	1	1	1	1	1	1	1	-	-
q10	1	1	1	1	1	1	1	1	1	1	1	-
q11	1	1	1	1	1	1	1	1	1	1	1	1

attention mask during third chunk prefill

Mask is a trapezoid matrix

Attention mask matrices for successive prefill chunks


Source: SARATHI: Efficient LLM Inference
by Piggybacking Decodes with Chunked
Prefills

Continuous Batching

- Continuous batching = ragged batching + dynamic batching

Continuous batching


Initial prompts: ["Who ate the pie", "I am sure this project", "The largest animal"]



Batch structure:
+ 1 prompt w/ 5 tokens (prefill)

Continuous Batching

- Continuous batching = ragged batching + dynamic batching



Batch structure:

+ 1 prompt w/ 1 tokens (decode)

+ 1 prompt w/ 4 tokens (chunked prefill)

Forward pass -----

Continuous Batching

- Continuous batching = ragged batching + dynamic batching



Batch structure:

- + 1 prompt w/ 3 tokens (end of chunked prefill)
- + 1 prompt w/ 2 tokens (chunked prefill)

Thanks!

