

LLM Inference Optimization

Autumn 2025

Lecturer: Yuedong (Steven) Xu

Shenzhen Loop Area Institute

yuedongxu@slai.edu.cn

Fudan University

ydxu@fudan.edu.cn

Disclaimer


Machine learning systems is a broad and rapidly evolving field. The course material has been developed using a broad spectrum of resources, including research papers, lecture slides, blogposts, research talks, tutorial videos, and other materials shared by the research community. Sometimes animations and pictures from external sources are heavily reused.

Inference Optimization: Outline

- Overview
- Attention Optimization
- Continuous Batching
- KV Cache Optimization
 - **Model Architecture Optimization**
 - KV Cache Data Structure Optimization
 - KV Cache Memory Optimization
- Speculative Decoding
- Distributed Serving

Attention Architecture Optimization

- Multi-Head Attention
 - Each head has its own Query, Key and Value matrices



Attention Architecture Optimization

- Multi-Query Attention (MQA)
 - All heads in a layer share a single set of W_K and W_V
 - Size of Q unchanged, and sizes of K and V shrink by h times


Component	Weight matrix shape	Number of parameters (exact)
Query projection	$W_Q \in \mathbb{R}^{d \times (h \cdot d_h)}$	$d \cdot (h \cdot d_h) = d^2$
Key projection	$W_K \in \mathbb{R}^{d \times d_h}$	$d \cdot d_h = d \cdot (d/h) = d^2/h$
Value projection	$W_V \in \mathbb{R}^{d \times d_h}$	$d \cdot d_h = d^2/h$
Output projection	$W_O \in \mathbb{R}^{(h \cdot d_h) \times d}$	$(h \cdot d_h) \cdot d = d^2$
Total	Shared	$2d^2 + 2 \cdot \frac{d^2}{h}$

$$\frac{MQA}{MHA} = \frac{1 + 1/h}{2}$$

Number of parameters in MQA


Attention Architecture Optimization

- Grouped Query Attention (GQA)
 - A group of heads share the same W_K and W_V




Attention Architecture Optimization

- Cross-Layer Attention
 - Consecutive layers (e.g. two) share the same K and V projections
 - Reduction in KV cache memory footprint (also applicable to model pre-training)
 - Orthogonal to MQA and GQA
 - Serious degradation in perplexity when # of shared layers exceeds two




Standard Attention



Cross-Layer Attention


Attention Architecture Optimization

- Decoupled-Head Attention
 - Insight: high similarity across some heads in a layer, and significant variation in the distribution of head-cluster similarities across layers




Attention Architecture Optimization

- Decoupled-Head Attention
 - Adaptively configures group sharing for key/value heads across various layers



Attention Architecture Optimization


- Multi-Head Latent Attention (MLA)
 - Latent vector: a small, shared representation that key and value are compressed into
 - Joint compression of keys and values



- During inference only $c_t^{KV} \in \mathbb{R}^{d_c}$ is **cached** where $d_c \ll d_h \cdot n_h$

Attention Architecture Optimization


- Multi-Head Latent Attention (MLA)
 - Reducing activations during *training*



- Decoupled Rotary Position Embedding (RoPE)
 - RoPE is position-sensitive, while the position information gets mixed into the compressed keys and values after using low-rank compression
 - Introducing additional multi-head queries q_t^R and a shared key k_t^R , and caching k_t^R

Attention Architecture Optimization

- Decoupled Rotary Position Embedding (RoPE)



The architecture of Multi-head Latent Attention (MLA)

$$\begin{aligned}
 & \text{cached} \\
 & q_t^R = \text{RoPE}(W^{QR} c_t^Q) \\
 & q_{t,i} = [q_{t,i}^C, q_{t,i}^R] \\
 & \boxed{k_t^R} = \text{RoPE}(W^{KR} h_t) \\
 & k_{t,i} = [k_{t,i}^C, k_{t,i}^R] \\
 & v_t^C = W^{UV} c_t^{KV} \\
 & q_{t,i} = \sum_{j=1}^t \text{Softmax}_j \left(\frac{q_{t,i}^T q_{j,i}^C}{\sqrt{d_h + d_h^R}} \right) v_{j,i}^C \\
 & u_t = W^O [o_{t,1}; o_{t,2}; \dots; o_{t,n_h}]
 \end{aligned}$$

Attention Architecture Optimization

- KV Cache Per-Token


MHA	GQA	MQA	MLA
$2d$	$2n_{kv}d_h$	$2d_h$	$d_c + d_h^{rope}$
32768 $(n_h = 128, d_h = 128)$	256~32768	256	576 $d_c = 512$

Inference Optimization: Outline

- Overview
- Attention Optimization
- Continuous Batching
- KV Cache Optimization
 - Model Architecture Optimization
 - **KV Cache Data Structure Optimization**
 - KV Cache Memory Optimization
- Speculative Decoding
- Distributed Serving

Memory Management


- Paged Attention
 - Uncertainty in **Decode**: no one knows how many tokens will be generated (remember continuous batching), and how much space needs to be allocated beforehand



- Memory fragmentations
 - Reserved fragmentation : reserved memory space not filled with output tokens temporarily
 - Internal fragmentation: reserved but never used even after the token generation completes
 - External fragmentation: memory space insufficient for inserting new KV cache

Memory Management

- Paged Attention
 - An example of memory fragmentation




NVIDIA A100 40GB

Grey: parameters persisting in GPU memory throughout serving

Red: KV cache (de)allocated per serving request

Yellow: used ephemeraly for activation.




NVIDIA A100 with 40GB HBM

Orca: another serving framework


Memory Management

- Paged Attention
 - Inspired by paged virtual memory
 - KV block table: recording the corresponding physical blocks of a logical block and the number of filled positions



Memory Management


- Paged Attention
 - Parallel sampling or beam search
 - Copy-on-Write (CoW): a memory optimization strategy where data isn't immediately duplicated when a copy is requested



Parallel sampling: logical blocks for the prompts of both sequences are mapped to the same physical blocks

Memory Management

- Revisiting PagedAttention
 - Requires re-writing the attention kernel and adds redundancy in the serving framework




A request allocates four KV cache blocks over time that are usually non-contiguous in virtual memory. The serving system needs to track the virtual memory addresses of KV cache blocks and pass them to the attention kernel at runtime.

- vAttention
 - Mitigates fragmentation in physical memory while retaining the virtual memory contiguity of the KV cache (supplementary information)

Memory Management


- Prefix Caching
 - 20-80 rule: a small portion of requests accounts for the majority of usage
 - KV cache reuse across requests
- Application Scenarios
 - Multi-Turn Conversations
 - Retrieval-Augmented Generation
 - Coding Assistants and IDE Integration
 - Few-Shot Prompting / In-Context Learning
 -




Around 50% of the tokens' KVCache in the real-world workloads can be reused

Memory Management


- KV Cache Reuse




(a) Few-shot learning



(b) Self-consistency



(c) Multi-turn chat




(d) Tree-of-thought

Memory Management

- Radix Tree

- Trie, a specialized search tree data structure used to store and retrieve strings
- Root does not store any character, each edge represents a character, and any path from the root to a node corresponds to a prefix of some string




A trie for keys "A", "to", "tea", "ted", "ten", "i", "in", and "inn". Each complete English word has an arbitrary integer value associated with it.

Memory Management

- Radix Tree


- A space-optimized trie (prefix tree) in which the number of children of every internal node is at most the radix r of this tree
 - Efficient for small sets and for sets of strings that share long prefixes
 - An illustration of storing five words and querying “toasting”




Memory Management

- Radix Tree


- Insert and delete




Insert 'water' at the root




Insert 'slower' while keeping 'slow'



Insert 'test' which is a prefix of 'tester'




Insert 'team' while splitting 'test' and creating a new edge label 'st'





Insert 'toast' while splitting 'te' and moving previous strings a level lower

Entire conversation cached




a new message from the first chat session



Radix Attention in SGLang

Memory Management

- Radix Attention
 - Cache-aware scheduling
 - For many requests in a queue, the order of executing them impacts the cache hit rate
 - For batch-processing, sort requests by matched prefix length and prioritize requests with longer matched prefixes instead of using a first-come first-served schedule
 - Compressed Finite State Machine
 - Adopted when tokens have strict formats, e.g. JSON and XML (supplementary information)




The decoding process of normal and compressed FSMs (the underscore _ means a space)

Memory Management

- Implementation Considerations
 - TreeNode structure
 - What to store: mapping to physical addresses, tracking node access patterns, recording the exact moment of a node lastly touched, pointers to parent or child nodes, etc.
 - Hash RadixAttention
 - Using hash value as the unique identification of a cache block
 - Node splitting and merging
 - The moments of updating RadixTree: request finished or a stage of chunked prefill finished
 - Eviction: memory constrained, time-to-leave expired, etc.
- Alternative Prefix Caching Frameworks
 - Prompt Cache, Shared Prefixes, etc.

Memory Management

- Radix Attention Eviction Policies
 - Classical approaches: Least recently used, Least frequently used, First-in-first-out, etc.
 - Machine learning based approaches




LPC: a MLP predictor and an eviction algorithm based on estimated conversation continuation probabilities

Inference Optimization: Outline



- Overview
- Attention Optimization
- Continuous Batching
- KV Cache Optimization
 - Model Architecture Optimization
 - KV Cache Data Structure Optimization
 - **KV Cache Memory Optimization**
- Speculative Decoding
- Distributed Serving

Memory Optimization




Memory Optimization

- Quantization
 - Fixed-precision or Mixed-precision Quantization, Outlier Redistribution
- Eviction
 - Permanent eviction: simple, but may cause an accuracy drop
 - Temporary eviction: offloads to CPU memory or SSD with limited bandwidth



KV Cache Offload: InfiniGen

- Key Insight
 - **Not** all KV cache data needs to stay in GPU memory at all times
 - **Limited** GPU high-bandwidth memory




Offloading to CPU memory is throttled by limited PCIe bandwidth

KV Cache Offload: InfiniGen


- Idea KV Cache Offloading

- Naïve offloading blocks the computation
- Prefetching large KV matrices cannot fully overlap communication and computation
- Only prefetching important KV pairs



A high-level illustration of the key idea

KV Cache Offload: InfiniGen




Q: Using attention pattern of layer i to estimate that of layer $i+1$?


A: NO! The tokens deemed unimportant in the current iteration could become important in subsequent iterations.

Original Attention: Layer i

Query Projection




Attention




Minimal Rehearsal: Layer $i - 1$

Partial Query Projection




Attention Speculation





Speculative KV prefetching: predicting the **important** tokens

Original Attention: Layer i




Minimal Rehearsal: Layer $i - 1$



Cosine Similarity

1.0
0.5
0.0



OPT-6.7B

OPT-13B

OPT-30B

Llama-2-7B

Llama-2-13B




in_{i-1}

$Attn_{i-1}$

FFN_{i-1}

Observation: attention inputs of consecutive attention layers are highly similar


Minimal Rehearsal: Layer $i - 1$



Before



×



=

Attention Score

33

9



×



=

24

Partial Query

Partial Key

After



×



=

Attention Score

33

1



×



=

32


Partial Query

Partial Key

Key/Query skewing using singular value decomposition


KV Cache Offload: InfiniGen

- Experimental results: OPT and LLama2 on NVIDIA RTX A6000 with PCIe 3.0 Gen




KV Cache Eviction: StreamingLLM

- Inference for Infinite Token Length
 - Insufficient GPU HBM, and high latency of KV cache offloading
 - Poor perplexity (PPL) for truncated attention methods (e.g. evicting old KV)



KV Cache Eviction: StreamingLLM



- Key Insights
 - Initial tokens have large attention scores, even if not semantically significant
 - **Attention sinks**: tokens that disproportionately attract attention irrespective of their relevance
 - Attributed to **Softmax** of autoregressive language modeling



Visualization of the average attention logits in Llama-2-7B over 256 sentences

KV Cache Eviction: StreamingLLM

- Serving Infinite Streams
 - Preserve the KV of attention sink tokens, keep the KV of tokens in the sliding window, and evict remaining tokens




An illustration of KV cache eviction


- Adjust the position of tokens. StreamingLLM focuses on positions within the cache rather than those in the original text

KV Cache Eviction: StreamingLLM

- Serving Infinite Streams




Accuracy on the StreamEval benchmark



Language modeling perplexity of StreamingLLM on super long texts with 4 million tokens





KV Cache for RAG: CacheBlend

- RAG: repetitive token sequences, but without common prefix




KV Cache for RAG: CacheBlend

- CacheBlend for RAG
 - Missing cross-attention (think why?)




KV Cache for RAG: CacheBlend


- CacheBlend for RAG
 - Selectively recomputing KV




Challenge and goal




(a) Default: Full KV re-compute.
Prefill on entire input



(b) Prior work: Prefix caching.
Only reusing *prefix's* KV cache



(c) Prior work: Full KV reuse.
Reusing all KV caches, ignoring cross-attention




(d) CacheBlend (ours): Selective KV re-compute.
Reusing all KV caches but re-computing a small fraction of KV

Contrasting full KV recompute, prefix caching, full KVreuse, and CacheBlend's selective KV recompute


KV Cache for RAG: CacheBlend

- Selectively recomputing KV
 - Key insight: ***sparsity*** of attention matrix and ***similarity*** of attention matrices on two consecutive layers
 - Attention deviation reduces as KV of more tokens on each layer are recomputed




KV Cache for RAG: CacheBlend


- CacheBlend for RAG
 - Selectively recomputing KV



Distribution of KV deviation of tokens on one layer




Rank correlation of KV deviation per token between two consecutive layers



CacheBlend selects high KV deviation tokens of one layer according to inter-layer dependency property

Supplementary Reading on Attention Sinks

- Do we need attention sinks? Voices from a different perspective
 - “Sparse gating mechanism mitigates ‘attention sink’ and enhances long-context extrapolation performance” – From Qwen



Inference Optimization: Outline

- Overview
- Attention Optimization
- Continuous Batching
- KV Cache Optimization
 - Model Architecture Optimization
 - KV Cache Management
 - KV Offloading and Transmission
- **Speculative Decoding**
- Distributed Serving

Speculative Decoding

- Autoregressive models decode one token at a time
 - Decoding is memory-bound—a factual observation rather than an assumption
- Some tokens are easier to generate than others
 - Using a smaller model to do the easy job

What is the square root of 7? The square root of 7 is 2.646.

Easy to generate
(i.e. with high probability)



Most uncertain




Speculative Decoding

- **Draft model:** a *smaller, lightweight* model that runs alongside to help speed up the main LLM's inference process
- **Target model:** an LLM verifies these proposed tokens *in parallel* and accepts those that match its own predictions
- Key assumption: draft model speculates a sequence of tokens very fast, target model accepts multiple tokens through one-round verification (inference).

Speculative Decoding

- Draft-then-Verify Paradigm



Speculative Decoding


- Verify In Parallel



The original LLM output: [BOS] This is good news. [EOS]

Speculative Decoding


- Verify In Parallel



Draft output coincides with original LLM by far

Speculative Decoding

- Verify In Parallel




Bifurcation: The index of the first draft token that fails verification.

Speculative Decoding

- Speculative decoding metrics
 - **Speculation accuracy** (high) and **drafting time** (low)
- Drafting
 - Self drafting
 - **Extra head**, layer skipping, mask predict
 - Independent drafting
 - **Non-Autoregressive LM**, **Small LM**, Context Retrieval, Copy
- Verification:
 - Strict matching, approximate matching


Speculative Decoding

- A brief history



Self-Drafting


- Block-wise Decoding (2018)
 - Augmenting the dimensionality of the final projection layer by a factor of k




- Apply the original vocabulary projection, add k output layers, add a hidden layer, compute k separate softmaxes per position
- Train the extra parameters or fine-tune them

Self-Drafting

- Block-wise Decoding (2018)
 - Predict-Verify-Accept versus Predict-Verify




Two forward passes



One forward pass

Self-Drafting

- Medusa (2018)



Self-Drafting

- Medusa (2018)

- Model training

- The original model stays untouched, and only the new heads are trained/fine-tuned

probability of token y predicted by the k -th head

Medusa-1: Frozen Backbone

$$\mathcal{L}_{\text{MEDUSA-1}} = \sum_{k=1}^K -\lambda_k \log p_t^{(k)}(y_{t+k+1}).$$

Medusa-2: Joint Training

$$\mathcal{L}_{\text{MEDUSA-2}} = \mathcal{L}_{\text{LM}} + \lambda_0 \mathcal{L}_{\text{MEDUSA-1}}$$

- Prediction


- Medusa heads predict multiple likely words for the corresponding position
 - The predicted tokens are combined and processed using a tree-based attention mechanism

- Verification

- A typical acceptance scheme is employed to pick the longest plausible prefix from the candidates for further decoding

Self-Drafting

- Medusa (2018)
 - Tree-based attention




- Using tree attention to process multiple candidates concurrently.
 - The top-2 predictions from the first MEDUSA head and the top-3 from the second result in a total of 6 candidates. Each of these candidates corresponds to a distinct branch within the tree structure.
- Devising an attention mask.
 - We exclusively permits attention flow from the current token back to its antecedent tokens.

Self-Drafting

- Draft & Verify (2023)

- Drafting Stage: selectively skipping certain intermediate layers
- Verification Stage: evaluating all drafted tokens in a full-model forward pass
- Unable to Draft and Verify in one pass



<s> Where is Zur ich ? Ans : Zur

<s> Where is Zur ich ? Ans : Zur **ich is the largest**

<s> Where is Zur ich ? Ans : Zur **ich is the largest city in Switzerland . It**

<s> Where is Zur ich ? Ans : Zur **ich is the largest city in Switzerland . It is located in**

<s> Where is Zur ich ? Ans : Zur **ich is the largest city in Switzerland . It is located in the eanton north**

<s> Where is Zur ich ? Ans : Zur **ich is the largest city in Switzerland . It is located in the north - central**

<s> Where is Zur ich ? Ans : Zur **ich is the largest city in Switzerland . It is located in the north - central Switzerland part**

<s> Where is Zur ich ? Ans : Zur **ich is the largest city in Switzerland . It is located in the north - central part of Switzerland. the**

<s> Where is Zur ich ? Ans : Zur **ich is the largest city in Switzerland . It is located in the north - central part of the country ...**

[token]: context tokens.


[token]: accepted draft tokens.

[token]: rejected draft tokens.

[token]: prediction from verification.


Self-Drafting

- Draft & Verify (2023)
 - Objective: the average inference time per verified token on a development set
 - Method: Using Bayesian optimization to find which layers to skip



Independent Drafting

- Non-autoregressive LMs: usually less performant




- Smaller Models for Drafting: tradeoff between model size and acceptance rate

Draft Model	Target Model	Acceptance Rate	Speed Up
T5-Small (77M)	T5-XXL (11B)	0.75	3.4x
T5-Base (250M)	T5-XXL (11B)	0.8	2.8x
T5-Large (800M)	T5-XXL (11B)	0.82	1.7x

Independent Drafting

- SpecInfer (2024)



Learning-based speculator: predict the LLM's output by maximizing the overlap between the speculated token tree and the tokens generated by the LLM using incremental decoding

Speculative Decoding: Brief Summary


- Key Idea
 - Employing small models to predict multiple tokens, and verify them via LLM
 - Requiring accurate prediction and low latency
- Feasibility
 - Some tasks are easy, while some are difficult
 - Decoding is usually not compute-bound, allowing fast computing of small models
- Methods
 - Drafting: self-drafting and independent drafting
 - Verification: speculative decoding (exact match) and speculative sampling (output-distribution match, and their versatile approaches)

Inference Optimization: Outline

- Overview
- Attention Optimization
- Continuous Batching
- KV Cache Optimization
 - Model Architecture Optimization
 - KV Cache Management
 - KV Offloading and Transmission
- Speculative Decoding
- **Distributed Serving**

Prefill-Decode Disaggregation


- Recalling continuous batching
 - Batching prefill and decoding phase together hurt both TTFT and TPOT
 - Three representative approaches: ***Orca***, ***Chunked Prefill*** and ***Prefill-Decode disaggregation***




Prefill and decode times are prolonged

Prefill-Decode Disaggregation

- Application-dependent SLO
 - **High throughput != high goodput**



LLM as a service: diverse SLO





Overprovisioning for SLO increases serving cost

Prefill-Decode Disaggregation


- Benefit of Disaggregation
 - An example

Disaggregation is a technique that



GPU

Max System rps
 $= \text{Min}(\text{Prefill}, \text{Decode})$
 $= 1.6 \text{ rps} / \text{GPU}$




GPU GPU GPU

Disaggregate (2P1D)
 $= \text{Min}(5.6 \times 2, 10) \text{ rps} / 3 \text{ GPU}$
 $= 3.3 \text{ rps} / \text{GPU}$


Prefill-Decode Disaggregation

- PD Disaggregation
 - Compute-bound prefill and memory-bound decode execute on different GPU pools
 - Prefill GPUs **transfer** KV cache to Decode GPUs
 - Optimizing prefill and decode separately via most suitable parallelism



Prefill-Decode Disaggregation

- Experimental Analysis





Evaluation of DistServe against vLLM on various datasets

Prefill-Decode Disaggregation

- Chunked Prefill versus Prefill-Decode Disaggregation
 - Memory footprint for model parameters
 - Co-located or disaggregated?
 - KV cache transfer on slow links
 - Co-located or disaggregated?
 - Token throughput prone applications
 - Co-located or disaggregated?
 - Heterogeneous GPUs
 - Co-located or disaggregated?


Prefill-Decode Disaggregation

- Chunked Prefill versus PD Disaggregation




Attention-FFN Disaggregation

- MoE Inference



Attention-FFN Disaggregation


- Motivation of AF Disaggregation
 - Attention is memory intensive and FFN is compute intensive in dense models
 - Sparse MoE models reduce the batch size on an expert



GPU utilization of Attention and FFN vs. batch size in dense model, MoE


Attention-FFN Disaggregation

- Key Idea of MegaScale-Infer
 - Independent scaling: feed more requests to FFN
 - Heterogeneous deployment: Adopt more cost-effective hardware for each module




Attention-FFN Disaggregation

- Challenges
 - Pipeline bubbles: sequential computation of a batch will result in pipeline bubbles, i.e. not all the GPUs being utilized at the same time
 - Ping-Pong pipeline parallelism: dividing a large batch into multiple microbatches




Attention-FFN Disaggregation

- Challenges
 - Simple idea, yet definitely non-trivial implementation
 - Minimize overhead such as data copies
 - CPU manages send/receive operations
 - GPU Direct over RDMA data transfer




Attention-FFN Disaggregation

- Experimental Analysis




Homogeneous clusters



Heterogeneous clusters

Memory Pooling for KV: MoonCake


- MoonCake: Trading more storage for less compute



Huge Challenges to Storage Systems


Memory Pooling for KV: MoonCake

- MoonCake
 - Distributed KV cache pool across all inference nodes
 - GPUDirect RDMA/Storage for high bandwidth
 - Prefix-hashed KV cache object storage
 - LRU eviction policy
 - Composite techniques including zero-copy, multi-NIC, topology-aware, load balancing, etc.




Memory Pooling for KV: MoonCake

- MoonCake
 - Effective request capacity and cache hit ratio



Number of requests that meet the latency requirements



Cache hit rate: global cache > local cache

Thanks!

